

# Let's Hunt a Memory Leak



**Sanket Patel**

Site Reliability Engineer @ LinkedIn

# Agenda

—

# Explore:

- Python and Objects
  - Allocs
  - Deallocs
- The Leak
  - with confirmed existence
- Hunting!

Won't Explore:

- Specifics of memory management and garbage collection
- Why did *Katappa* kill *Bahubali*

# Python 2 and Integers

---

```
Python 2.7.15 (default, Jan 12 2019, 21:07:57)
[GCC 4.2.1 Compatible Apple LLVM 10.0.0 (clang-1000.11.45.5)] on darwin
Type "help", "copyright", "credits" or "license" for more information.

>>> import os, psutil, gc, time

>>> l=[i for i in range(100000000)]

>>> print(psutil.Process(os.getpid()).memory_info())
pmem(rss=3244871680L, vms=7824240640L, pfaults=1365384, pageins=460)
```

# Python 2 and Integers

---

```
Python 2.7.15 (default, Jan 12 2019, 21:07:57)
[GCC 4.2.1 Compatible Apple LLVM 10.0.0 (clang-1000.11.45.5)] on darwin
Type "help", "copyright", "credits" or "license" for more information.

>>> import os, psutil, gc, time

>>> l=[i for i in range(100000000)]

>>> print(psutil.Process(os.getpid()).memory_info())
pmem(rss=3244871680L, vms=7824240640L, pfaults=1365384, pageins=460)

>>> del l
```

# Python 2 and Integers

---

```
Python 2.7.15 (default, Jan 12 2019, 21:07:57)
[GCC 4.2.1 Compatible Apple LLVM 10.0.0 (clang-1000.11.45.5)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>> import os, psutil, gc, time

>>> l=[i for i in range(100000000)]
>>> print(psutil.Process(os.getpid()).memory_info())
pmem(rss=3244871680L, vms=7824240640L, pfaulsts=1365384, pageins=460)

>>> del l
>>> print(psutil.Process(os.getpid()).memory_info())
pmem(rss=2509352960L, vms=6964514816L, pfaulsts=1381131, pageins=460)
```



# Objects

---



**THEY SAY**

Everything in Python is an Object

---

# What is it?

```
typedef struct _object {  
    PyObject_HEAD      // contains ref_count  
} PyObject;
```

# What is it?

```
typedef struct _object {  
    PyObject_HEAD      // contains ref_count  
} PyObject;
```

```
typedef struct {  
    PyObject_HEAD  
    long ob_ival;  
} PyIntObject;
```

# Free Lists

---

```
static PyListObject *free_list[PyList_MAXFREELIST];  
static PyDictObject *free_list[PyDict_MAXFREELIST];
```

# Free Lists

---

```
static PyListObject *free_list[PyList_MAXFREELIST];
static PyDictObject *free_list[PyDict_MAXFREELIST];

static PyFrameObject *free_list = NULL;
static int numfree = 0;          /* number of frames currently in free_list */
#define PyFrame_MAXFREELIST 200
```

# Garbage Collection

---

- Generation based
  - A linked list for each generation
- When run in last generation (= 2)
  - Clear free lists as well

```
/* Clear free list only during the collection of the highest
 * generation */
if (generation == NUM_GENERATIONS-1) {
    clear_freelists();
}
```



# Allocation

ref\_count++



## 1. If free\_list has space

- Use last slot from free\_list

## 2. Else: allocate memory

- Initialize the object
- Register object with GC (ie: add in linked list)\*



not every object gets GC tracked

---

## 1. Immutables

- long, float, string

## 2. containers

- list, dict

# Deallocation

when `ref_count` drops to 0

---

1. Untrack from GC (ie: modify pointers)\*
2. Decrement ref count of contained types\*
3. If: `free_list` has space
  - append to `free_list`
4. Else: `free_list` is full
  - free the memory

# Recap

---

1. free\_list
2. Generational GC
3. Alloc
  - ref\_cnt++
  - free\_list or alloc
  - GC tracking\*
4. Dealloc
  - ref\_cnt == 0
  - GC untrack
  - free\_list or dealloc

# Python 2 and Integers

---

free\_list: **UNLIMITED** IN SIZE

Why?

- SPEED!
- Allocate space for N(=24) objects in one go.
- Deallocate never (except GC in gen 2)

<https://superuser.blog/python-2-integers/>

# The Leak

—



# A Good 'ol Flask App

---

## What does it do?

- Serve metrics for services

## What does it contain?

- Fake metrics generator
- Helper which converts response in JSON
- Performance profiling
- An amplified leak!

# The App

---

```
from flask import Flask
from helpers import _get_service_metrics, json_response

app = Flask('pycon')

@app.route('/metrics/<service>')
@json_response
def get_service_metrics(service):
    # returns list of metric data points
    data = _get_service_metrics(service)
    return data
```

# Memory Usage

---

## Initial Memory Usage:

```
~/workspace/pycon $ ./mem_usage.sh app.py  
24MB
```

## Sending *Some* Traffic:

```
/workspace/pycon $ ab -n50 http://127.0.0.1:5000/metrics/test_service
```

## Memory Usage After:

```
~/workspace/pycon $ ./mem_usage.sh app.py  
4085MB
```

# The App

with added GC

---

```
import gc
from flask import Flask
from helpers import _get_service_metrics, json_response

app = Flask('pycon')

@app.route('/metrics/<service>')
@json_response
def get_service_metrics(service):
    # returns list of metric data points
    data = _get_service_metrics(service)
    gc.collect()
    return data
```

# Memory Usage

with added GC

---

Initial Memory Usage:

```
~/workspace/pycon $ ./mem_usage.sh app.py  
24MB
```

Sending *Some* Traffic:

```
/workspace/pycon $ ab -n50 http://127.0.0.1:5000/metrics/test_service
```

Memory Usage After:

```
~/workspace/pycon $ ./mem_usage.sh app.py  
4070MB
```

There's a leak!





# Let's Hunt!

---

# Tools Available

---

## Core dump analysis

### Native libraries:

- objgraph
- memory-profiler
- Heapy
- tracemalloc
- ...

# tracemalloc

because built-in, that's why!



# tracemalloc

straight from the docs

---

- Traceback where an object was allocated
- Statistics on allocated memory blocks **per filename and per line number.**
- Compute the differences between two snapshots **to detect memory leaks**

# tracemalloc

how to?

---

```
import tracemalloc

tracemalloc.start()

# ... do stuff ...

snapshot = tracemalloc.take_snapshot()
top_stats = snapshot.statistics('lineno')

print("[ Top 10 ]")
for stat in top_stats[:10]:
    print(stat)
```

# Inside tracemalloc

and the magic it does!

---

## Python Memory Operations API

- PyObject\_Malloc/Free
- PyMem\_Malloc/Free

Override them.

Trace the calls.



# The App

again!

---

- Stateless
- Fetches metrics and serves them

# The Master Plan

to catch the Leak

---

1. Warm up
  - by sending a couple requests
2. Snapshot
3. Send traffic
4. Snapshot again
5. Analysis

# The Plan Executor

---

```
import tracemalloc
tracemalloc.start()
s=None

...
def get_service_metrics(service):
...

@app.route("/snapshot")
def snap():
    global s
    if not s:
        s = tracemalloc.take_snapshot()
        return "taken snapshot\n"
    else:
        lines = []
        top_stats = tracemalloc.take_snapshot().compare_to(s, 'lineno')
        for stat in top_stats[:5]:
            lines.append(str(stat))
        return "\n".join(lines)
```

# Execution

---

Warm up:

```
~/workspace/pycon $ ab -n2 http://127.0.0.1:5000/metrics/test_service  
~/workspace/pycon $ curl 127.0.0.1:5000/snapshot  
taken snapshot
```

Sending *Some* Traffic:

```
/workspace/pycon $ ab -n50 http://127.0.0.1:5000/metrics/test_service
```

# **\*drumrolls\***

---

The leak:

```
~/workspace/pycon $ curl 127.0.0.1:5000/snapshot
workspace/pycon/helpers.py:25: size=4206 MiB (+4034 MiB), count=1 (+0), average=4206 MiB
/usr/local/lib/python3.7/site-packages/werkzeug/serving.py:226: size=21.7 KiB (+18.5 KiB),
count=20 (+17), average=1112 B
/usr/local/Cellar/python/3.7.4/Frameworks/Python.framework/Versions/3.7/lib/python3.7/socket
.py:253: size=14.4 KiB (+5168 B), count=42 (+34), average=350 B
/usr/local/lib/python3.7/site-packages/werkzeug/_compat.py:193: size=4128 B (+3534 B),
count=60 (+51), average=69 B
/usr/local/lib/python3.7/site-packages/werkzeug/serving.py:223: size=3640 B (+3094 B),
count=60 (+51), average=61 B
```

# The Culprit

---

```
from flask import jsonify

PERF_DATA = dict()

def json_response(func):
    def wrapper(*args, **kwargs):
        global PERF_DATA
        start = time.time()
        r_val = func(*args, **kwargs)
        time_taken = start - time.time()

        # amplify the leak
        time_taken = [time_taken]*10000000
        # update performance data
        key = str((args, kwargs))
        if key in PERF_DATA:
            PERF_DATA[key].extend(time_taken)
        else:
            PERF_DATA[key] = time_taken
        return jsonify(r_val)

    return wrapper
```

# Recap

---

We

- Understood the memory management
- Understood the application behaviour
- Hunted the leak!



# That's all for today...

My home is @ <https://sanket.plus>

